

Model Overview

Name:
nvidia/Audio2Emotion
(v2.2 production,
v3.0 experimental)

Task:
Speech Emotion
Recognition (SER)

Architecture:
Wav2Vec2 backbone
+ SER head

6 output classes:
anger, disgust
fear, joy
neutral, sadness


Input:
Mono PCM 16-bit
Any SR (->16 kHz)

Output:
Per-frame emotion
scores [0,1]

Context:
Audio2Face-3D
pipeline component
Drives 3D avatar
facial expressions

NVIDIA Audio2Emotion – Inference Pipeline

nvidia/Audio2Emotion-v2.2 (v3.0) | Wav2Vec2-based SER | Part of Audio2Face-3D stack



Audio2Face-3D full pipeline (C++/CUDA SDK – faster than real-time, 60+ FPS)