

Pipeline Overview

Input: any audio
(wav / mp3 / flac)

Output: RTTM
“who spoke when”

3 core models:

1. segmentation-3.0
PyanNet
SincNet + Bi-LSTM
Powerset 7-class
2. wespeaker
resnet34-LM
ResNet34 encoder
LM fine-tuned
3. speaker-diarization
3.1 pipeline
Agglom. clustering

Key params:

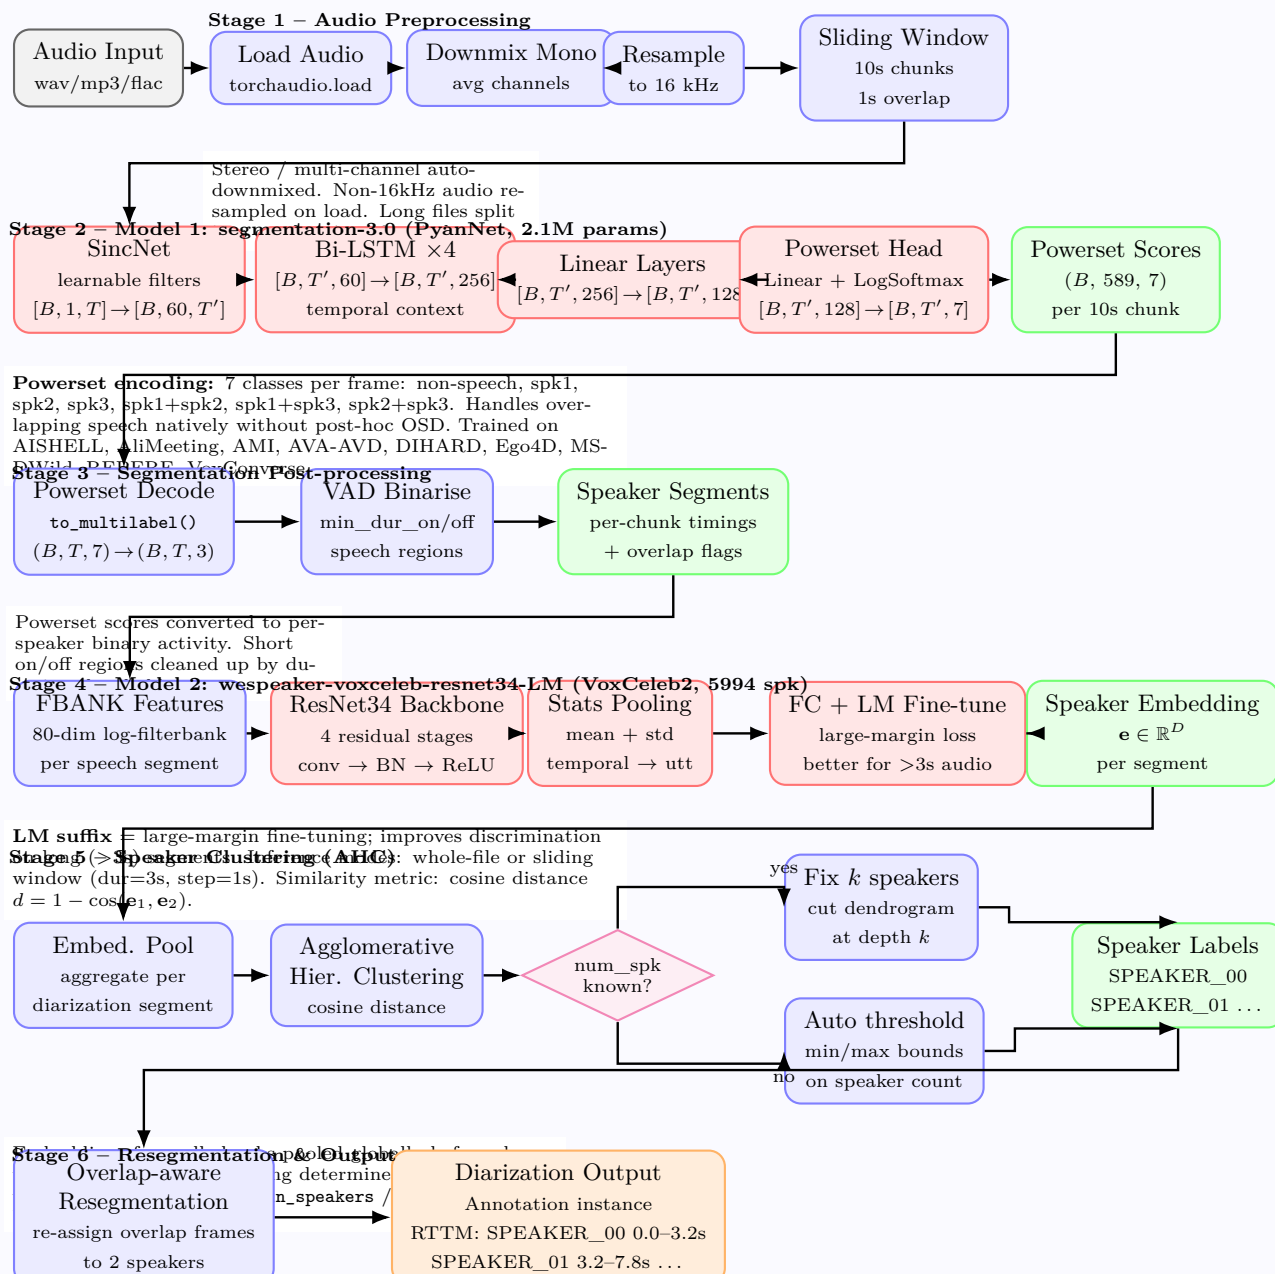
num_speakers
min/max_speakers
cluster threshold

Constraints:

Mono, 16 kHz
Pure PyTorch
pyannote.audio 3.1+

pyannote Speaker Diarization – Full Inference Pipeline

segmentation-3.0 | wespeaker-voxceleb-resnet34-LM | speaker-diarization-3.1



speaker-diarization-3.1 (pure PyTorch, pyannote.audio ≥ 3.1)